

COVID-19: Hydroxychloroquine Experiment Result Expansion

Patrick Poleshuk

8/20/2021

```
set.seed(101)
# Bayesian regression (using the R package 'rstanarm') uses Markov-chain Monte Carlo
# simulation to supply the parameter estimates. Essentially, to find our regression
# coefficients, we need to find the posterior distribution that they lie on, a probability
# distribution that is proportional to the maximum likelihood estimation of our coefficients
# and our prior beliefs of the probability distribution that our parameter of interest follows.
# The algorithm is complex, but essentially we can construct a random sample from our
# posterior distribution, where each estimate sampled is influenced by the one that came
# before it (called a Markov Chain). Using Markov chain sampling, we can draw samples from
# the posterior distribution without knowing the complete properties of said distribution. Under
# standard Monte Carlo simulation, samples (which follow a distribution) are made independent,
# yet using a Markov chain, in this case, allows us to hone in on the posterior distribution
# where we believe our true coefficient of effect lies.

# 4,000 samples are taken into account, and the median values of these probabilities
# are our best estimate for beta (which is assumed to follow a probability
# distribution in Bayesian statistics, rather than simply being a fixed, point-estimate as in
# the Frequentist case). To reiterate, our estimates derive from generating random values
# from the posterior PDF, where our best estimate beta is stationed,
# so it stands that every time we run a bayesian regression our best estimate will change
# very slightly since our random samples will fluctuate. The change isn't significant; however, we s
#

# Loading up required libraries & the excel data set of the experiment outcomes.
library(rstanarm)
library(dplyr)
library(tidyverse)

gautret <- readxl::read_excel('/Users/patrickpoleshuk/gautretData_2.xlsx')
names(gautret)[1] <- c('age')

# Before conducting any kind of regression analysis, I will like to visualize any
# structural differences between the control and treatment groups, in regards to
# age, sex, and disease classification.


g = gautret %>% mutate(category = ifelse(ctrCode == 0, "Control",
ifelse(ctrCode == 1 & atrCode == 0, "Hydrox", "Hydrox/Azithro")))

old <- theme_set(theme_bw())
# Setting a theme for ggplot.
```

```

dat = as.data.frame(g %>% group_by(category) %>% summarise(avg = mean(age)))
ggplot(data = g, aes(x = category, y = age, fill = category)) +
  geom_boxplot(outlier.shape = NA) + geom_jitter(width = 0.2) +
  geom_point(data = dat, aes(x = category, y = avg), color = "purple", shape = 8,
    size = 4) + labs(x = "Assigned Category",
    y = "Age",
title = "Box Plot of The Age Distribution by Group Assignment") +
  scale_fill_discrete(name = 'Assignment')

```



We can see from the box plot that the median, mean, & deviation from the mean, regarding age, differs between the control and treatment groups.

For a numerical representation:

```

stat_age = g %>% group_by(category) %>% summarise(Average = mean(age),
  Median = median(age),
  StDev = sd(age),
  IQR = IQR(age))

```

Visualizing dose concentration values:

```

g$conc = as.numeric(g$conc)
g$conc[is.na(g$conc)] = 0

```

```

data = g %>% group_by(category)


```

```

dat = as.data.frame(g %>% group_by(category) %>% summarise(avg = mean(conc)))

```

```
ggplot(data = g, aes(x = category, y = conc, fill = category)) +
  geom_boxplot() + geom_jitter(width = .2) +
  labs(x = "Assigned Category",
       y = "Concentration of Hydroxychloroquine ( $\mu$  g/ml)",
       title = "Box Plot of The Treatment Concentration by Assignment") +
  scale_fill_discrete(name = 'Assignment') +
  geom_point(data = dat, aes(x = category, y = avg), color = "purple", shape = 8,
            size = 4)
```



```
stat_conc = g %>% group_by(category)%>% summarise(Average = mean(conc),
                                                  Median = median(conc),
                                                  StDev = sd(conc),
                                                  IQR = IQR(conc))
```

*# For those exposed to either just the hydroxychloroquine treatment
 # or hydroxychloroquine paired with the azithromycin, we can get a good
 # idea of the concentration of the treatment within the bodies of the
 # subjects. We can see that, even though all subjects in the treatment
 # were dosed with 600 micrograms per milliliter, a higher dose resonates
 # for those who supplemented with azithromycin. Whether this is the result
 # of some differing physiological trait that either enhances or hinders the
 # the effect of the dosage on bodily drug concentration, a synergistic
 # impact of the azithromycin, or simply the result of an experimental
 # blunder in distributing the dosages, I am not sure.*

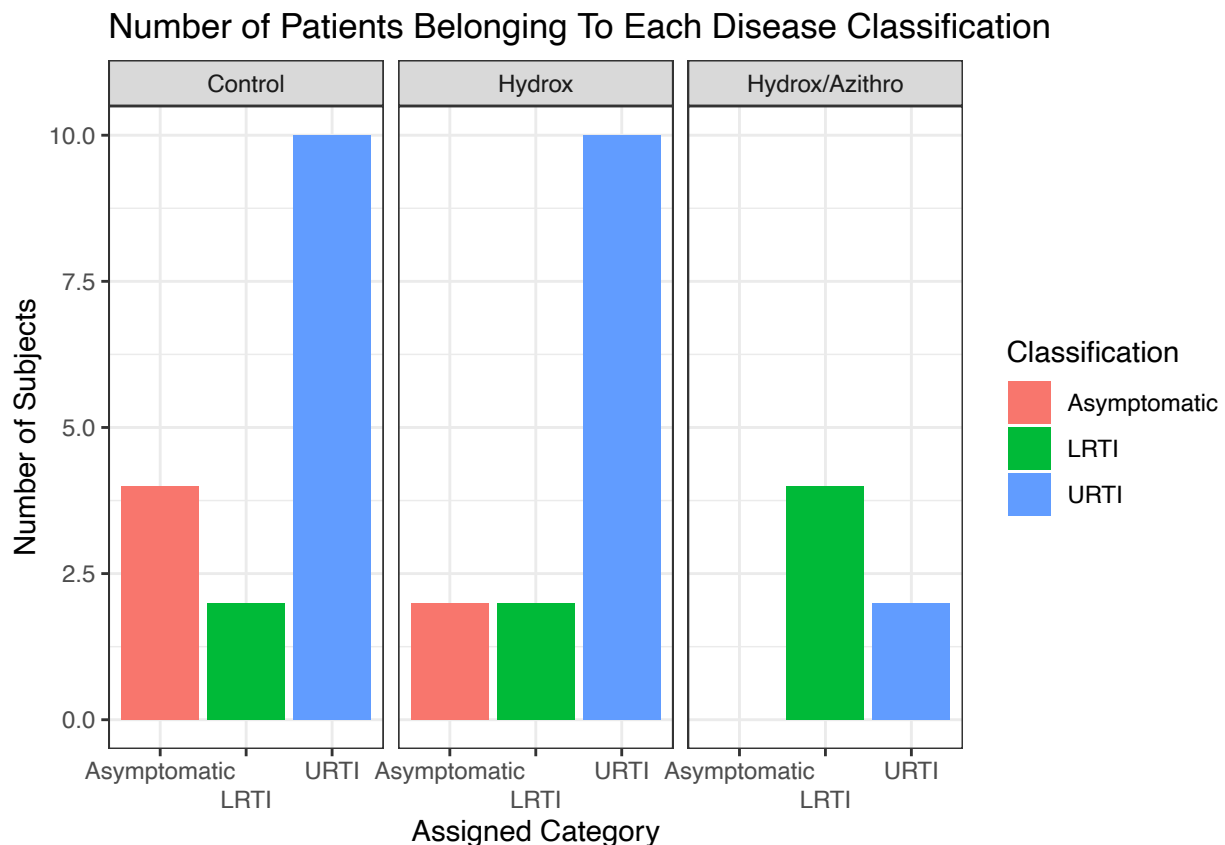
```

# Onto Disease classification differences:
g[31, 3] = 'LRTI'
# Correcting a data entry discrepancy in the excel data spreadsheet.

bar = g %>% group_by(category, status) %>%
  summarise(len = length(status))

library(scales)
ggplot(data = bar, aes(x = status, y = len, fill = as.factor(status))) +
  geom_col() +
  facet_wrap(~category) + scale_x_discrete(guide = guide_axis(n.dodge = 2)) +
  scale_fill_discrete(name = 'Classification') +
  labs(x = 'Assigned Category',
       y = 'Number of Subjects',
       title = "Number of Patients Belonging To Each Disease Classification")

```




The groups remained relatively balanced in regards to the number of URTI/LRTI/Asymptomatic patients contained within them. Still, we see no significant systematic difference that would bias the findings in the clinical experiment.

```

# Sex differences:
table = g %>% group_by(sex, category) %>% summarise(n = n())
ggplot(data = table, aes(x = category, y = n, fill = as.factor(category))) +
  geom_col() +
  facet_wrap(~sex) + scale_x_discrete(guide = guide_axis(n.dodge = 2)) +

```

```
scale_fill_discrete(name = 'Status') + labs(x = 'Faceted By Sex (Female/Male)',
      y = 'Number of Patients Per Assignment',
      title = 'Count of Grouping Assignment by Sex (Female & Male)')
```



*# Roughly an equally balanced amount of females and males in each group, when accounting for
the fact that there were more female patients studied in this experiment.*

```
set.seed(101)
```

*# Even though we have observed that selection bias should be minimal, with regards to
confounding our hydroxychloroquine & azithromycin treatment effect on COVID-19 test
results, given that all three assignment groups share a similar set of characteristics, we
can still examine whether these characteristics influence the outcome of interest.*

Regression Models:

t6 ~ Age

```
logistic_age = stan_glm(data = g, t6 ~ age, family = binomial(link = 'logit'))
```

```
##
```

```
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 1).
```

```
## Chain 1:
```

```
## Chain 1: Gradient evaluation took 0.000879 seconds
```

```
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 8.79 seconds.
```

```
## Chain 1: Adjust your expectations accordingly!
```

```

## Chain 1:
## Chain 1:
## Chain 1: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 1: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 1: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 1: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 1: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 1: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 1: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 1: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 1: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 1: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 1: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.061015 seconds (Warm-up)
## Chain 1:                0.058248 seconds (Sampling)
## Chain 1:                0.119263 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 1.8e-05 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.18 seconds.
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 2: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 2: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 2: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 2: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 2: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 2: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 2: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 2: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 2: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 2: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 2: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 0.057672 seconds (Warm-up)
## Chain 2:                0.056904 seconds (Sampling)
## Chain 2:                0.114576 seconds (Total)
## Chain 2:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 3).
## Chain 3:
## Chain 3: Gradient evaluation took 1.4e-05 seconds
## Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.14 seconds.
## Chain 3: Adjust your expectations accordingly!
## Chain 3:
## Chain 3:
## Chain 3: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 3: Iteration:   200 / 2000 [ 10%] (Warmup)

```

```

## Chain 3: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 3: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 3: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 3: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 3: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 3: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 3: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 3: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 3: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 3: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 3:
## Chain 3: Elapsed Time: 0.054374 seconds (Warm-up)
## Chain 3: 0.076701 seconds (Sampling)
## Chain 3: 0.131075 seconds (Total)
## Chain 3:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 4).
## Chain 4:
## Chain 4: Gradient evaluation took 2.1e-05 seconds
## Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.21 seconds.
## Chain 4: Adjust your expectations accordingly!
## Chain 4:
## Chain 4:
## Chain 4: Iteration: 1 / 2000 [ 0%] (Warmup)
## Chain 4: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 4: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 4: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 4: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 4: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 4: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 4: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 4: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 4: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 4: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 4: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 4:
## Chain 4: Elapsed Time: 0.0722 seconds (Warm-up)
## Chain 4: 0.070193 seconds (Sampling)
## Chain 4: 0.142393 seconds (Total)
## Chain 4:

```

```
print(logistic_age, digits = 4)
```

```

## stan_glm
## family: binomial [logit]
## formula: t6 ~ age
## observations: 36
## predictors: 2
## -----
##           Median MAD_SD
## (Intercept) 0.2261 0.7714
## age         0.0027 0.0155
##
## -----

```

```
## * For help interpreting the printed output see ?print.stanreg
## * For info on the priors used see ?prior_summary.stanreg

# Given that our dependent variable follows a Bernoulli distribution (our random variable
# can take only 2 possible outcomes as either you test negative or positive for COVID-19),
# we can estimate using a logistic regression. This allows us to find the probability that our
# dependent variable is true, given what we observe from our independent variables.

# More specifically, the coefficients we receive will tell us the log-odds that we test
# positive, given a 1 unit increase in our continuous independent variable, age.
# To give a better interpretation, we can standardize our age variable and take the inverse
# logit of the coefficient to convert our log-odds value into a probability. As it is now, our
# intercept value tells us that our model predicts that the log-odds of contracting COVID
# when someone is of age 0 is 0.2381. Additionally, for every one year increase in age, the log
# odds of contracting COVID increase by 0.0024. As we can see, this doesn't make for a very
# useful interpretation.

inv_logit <- function(x) {
  formula = exp(x) / (1 + exp(x))
  return(formula)
}
g$agesc = (g$age - mean(g$age))/(sd(g$age))
logistic_age_standardized = stan_glm(data = g, t6 ~ agesc,
                                     family = binomial(link = 'logit'))

##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 2.5e-05 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.25 seconds.
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: Iteration:    1 / 2000 [ 0%] (Warmup)
## Chain 1: Iteration:  200 / 2000 [ 10%] (Warmup)
## Chain 1: Iteration:  400 / 2000 [ 20%] (Warmup)
## Chain 1: Iteration:  600 / 2000 [ 30%] (Warmup)
## Chain 1: Iteration:  800 / 2000 [ 40%] (Warmup)
## Chain 1: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 1: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 1: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 1: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 1: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 1: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.054726 seconds (Warm-up)
## Chain 1:                0.057627 seconds (Sampling)
## Chain 1:                0.112353 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 2).
## Chain 2:
```



```

## Chain 2: Gradient evaluation took 1.6e-05 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.16 seconds.
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 2: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 2: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 2: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 2: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 2: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 2: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 2: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 2: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 2: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 2: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 2: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 0.061444 seconds (Warm-up)
## Chain 2:                    0.059849 seconds (Sampling)
## Chain 2:                    0.121293 seconds (Total)
## Chain 2:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 3).
## Chain 3:
## Chain 3: Gradient evaluation took 1.7e-05 seconds
## Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.17 seconds.
## Chain 3: Adjust your expectations accordingly!
## Chain 3:
## Chain 3:
## Chain 3: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 3: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 3: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 3: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 3: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 3: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 3: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 3: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 3: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 3: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 3: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 3: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 3:
## Chain 3: Elapsed Time: 0.054556 seconds (Warm-up)
## Chain 3:                    0.055586 seconds (Sampling)
## Chain 3:                    0.110142 seconds (Total)
## Chain 3:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 4).
## Chain 4:
## Chain 4: Gradient evaluation took 1.6e-05 seconds
## Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.16 seconds.
## Chain 4: Adjust your expectations accordingly!
## Chain 4:

```

```
## Chain 4:
## Chain 4: Iteration: 1 / 2000 [ 0%] (Warmup)
## Chain 4: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 4: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 4: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 4: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 4: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 4: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 4: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 4: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 4: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 4: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 4: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 4:
## Chain 4: Elapsed Time: 0.061886 seconds (Warm-up)
## Chain 4: 0.055614 seconds (Sampling)
## Chain 4: 0.1175 seconds (Total)
## Chain 4:
```

```
print(logistic_age_standardized, digits = 4)
```

```
## stan_glm
## family: binomial [logit]
## formula: t6 ~ agesc
## observations: 36
## predictors: 2
## -----
## Median MAD_SD
## (Intercept) 0.3574 0.3485
## agesc 0.0486 0.3468
## -----
## * For help interpreting the printed output see ?print.stanreg
## * For info on the priors used see ?prior_summary.stanreg
```

```
print(inv_logit(logistic_age_standardized$coefficients))
```

```
## (Intercept) agesc
## 0.5884221 0.5121460
```

Here we can see that our intercept value tells us the probability of one testing positive for COVID, given that they are of average age: 0.5884221. Our beta 1 coefficient gives us insight into the probability of testing positive for COVID for every 1 standard deviation we move away from the mean age observed. To give an example, where someone is 1 standard deviation away from the mean age:

```
int = logistic_age_standardized$coefficients[1]
beta1 = logistic_age_standardized$coefficients[2]
inv_logit(int + 1*beta1)
```

```
## (Intercept)
## 0.6001378
```

*# We see that someone who fits this description is predicted to have a 0.6001378 probability
of testing positive.*

```
inv_logit(int + 2.5*beta1)
```

```
## (Intercept)
## 0.6174951
```

*# Meanwhile, someone who is, say, 2.5 standard deviations away from the mean age is predicted
to have a 0.6174951 probability of testing positive for COVID.*

*# In other words, we can see that older people are more susceptible to test positive for COVID,
given our observed data, however, the effect is negligible.*

```
posterior_interval(logistic_age_standardized)
```

```
##              5%          95%
## (Intercept) -0.2166741 0.9551938
## agesc       -0.5115072 0.6569168
```

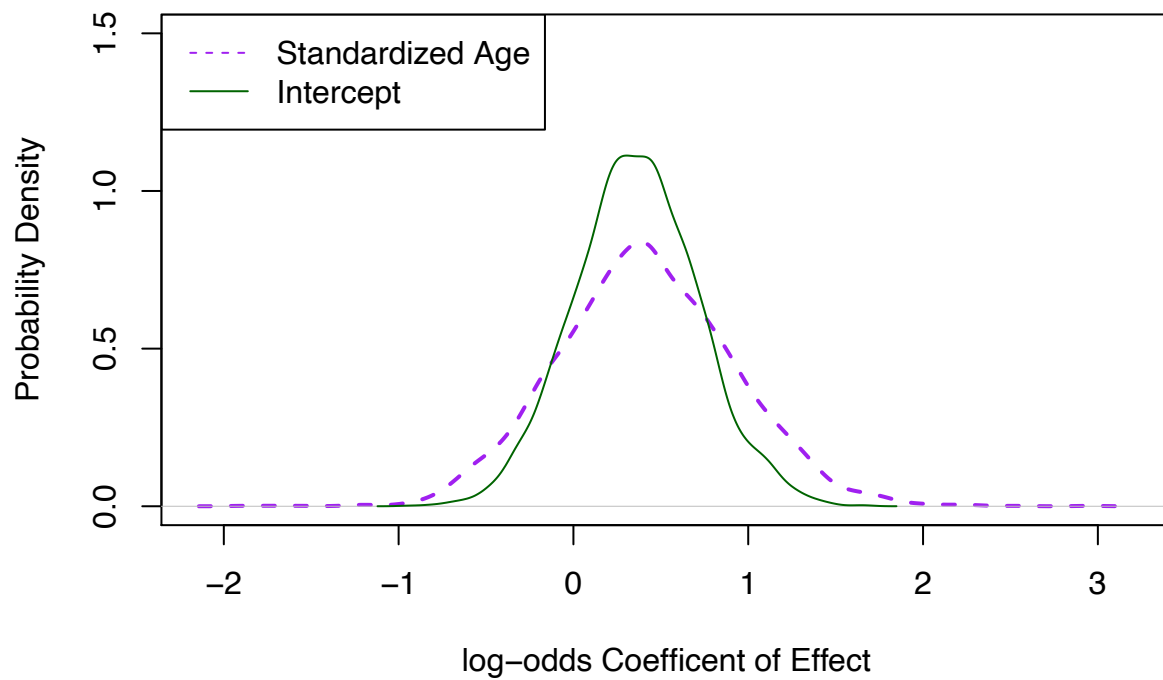
*# Unfortunately and predictably, our Bayesian credible interval tells us that our estimator
is inefficient, given that 0 is a plausible value for both these estimators. Our
interpretation for our beta 1 estimate, thus, is that there is a 95% probability that the
log-odds effect, for every 1 standard deviation shift away from the average mean observed,
on testing positive is between -0.5115072 & 0.6569168.*

*# To illustrate this with the posterior distribution of our parameter estimates: we can see
that the distribution of our parameters vary way too much for us to determine that this
effect of age, in this study, provided a meaningful effect on the susceptibility of testing
positive for COVID.*

```
sims = as.matrix(logistic_age_standardized)
int = sims[, 1]
agesc = int + sims[, 2]
```

```
plot(density(agesc), col = 'purple', lty = 2, lwd = 2,
     xlab = "log-odds Coefficient of Effect",
     main = "Posterior Density Function: Logistic Age Model",
     y = "Probability Density", ylim = c(0, 1.5))
lines(density(int), lwd = 1, col = 'darkgreen')
legend('topleft', legend = c("Standardized Age", "Intercept"), col = c("purple", "darkgreen"), lty =
```

Posterior Density Function: Logistic Age Model



```
sequence = seq(min(g$age)-1, max(g$age)+1, .01)

sims = as.matrix(logistic_age)


int = sims[, 1]
age = sims[, 2]

lower = c()
higher = c()
med = c()

for (i in 1:length(sequence)) {
  ypred = int + age*sequence[i]
  lower[i] = quantile(ypred, prob = .025)
  higher[i] = quantile(ypred, prob = .975)
  med[i] = median(ypred)
}

plot(g$age, g$t6, pch = 16, xlab = "Age of Subject",
     ylab = "Probability of Testing Positive for COVID-19",
     main = "Logisitic Regression of COVID Testing Outcome Regressed on Age")
lines(sequence, inv_logit(lower), lty = 2, col = "cyan")
lines(sequence, inv_logit(higher), lty = 2, col = "cyan")
lines(sequence, inv_logit(med), lwd = 2, col = "magenta")
```

Logisitic Regression of COVID Testing Outcome Regressed on Age



As we can see visually, given that age lacks the significance of determining a robust probabilistic outcome for COVID-19 test results, our logit model doesn't even remotely resemble a sigmoid curve. I should reiterate that this is desirable, in this case, as age could very well have been a meaningful confounder. Thankfully, however, that is not shown to be the case here.

For variables like this, plotting the logistic regression would be ineffective, and it would be better to just visualize the posterior density function shown above.

```
# t6 ~ Sex
logistic_sex = stan_glm(data = g, t6 ~ sexCode, family = binomial(link = 'logit'))
```

```
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 2e-05 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.2 seconds.
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 1: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 1: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 1: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 1: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 1: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 1: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 1: Iteration: 1200 / 2000 [ 60%] (Sampling)
```

```

## Chain 1: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 1: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 1: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.054064 seconds (Warm-up)
## Chain 1: 0.052196 seconds (Sampling)
## Chain 1: 0.10626 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 2.3e-05 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.23 seconds.
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration: 1 / 2000 [ 0%] (Warmup)
## Chain 2: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 2: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 2: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 2: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 2: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 2: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 2: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 2: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 2: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 2: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 2: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 0.051393 seconds (Warm-up)
## Chain 2: 0.057781 seconds (Sampling)
## Chain 2: 0.109174 seconds (Total)
## Chain 2:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 3).
## Chain 3:
## Chain 3: Gradient evaluation took 1.4e-05 seconds
## Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.14 seconds.
## Chain 3: Adjust your expectations accordingly!
## Chain 3:
## Chain 3:
## Chain 3: Iteration: 1 / 2000 [ 0%] (Warmup)
## Chain 3: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 3: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 3: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 3: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 3: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 3: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 3: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 3: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 3: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 3: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 3: Iteration: 2000 / 2000 [100%] (Sampling)

```

```

## Chain 3:
## Chain 3: Elapsed Time: 0.04998 seconds (Warm-up)
## Chain 3:           0.057465 seconds (Sampling)
## Chain 3:           0.107445 seconds (Total)
## Chain 3:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 4).
## Chain 4:
## Chain 4: Gradient evaluation took 1.8e-05 seconds
## Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.18 seconds.
## Chain 4: Adjust your expectations accordingly!
## Chain 4:
## Chain 4:
## Chain 4: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 4: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 4: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 4: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 4: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 4: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 4: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 4: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 4: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 4: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 4: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 4: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 4:
## Chain 4: Elapsed Time: 0.047397 seconds (Warm-up)
## Chain 4:           0.054426 seconds (Sampling)
## Chain 4:           0.101823 seconds (Total)
## Chain 4:

```

```
inv_logit(logistic_sex$coefficients)
```

```

## (Intercept)      sexCode
##   0.4719344    0.6884330

```

```

# If one is male, the probability that they test positive is 0.4733507. If one is female,
# the probability increases to 0.6903189.

```

```

# log-odds Estimates:
print(logistic_sex, digits = 4)

```

```

## stan_glm
## family:      binomial [logit]
## formula:     t6 ~ sexCode
## observations: 36
## predictors:  2
## -----
##           Median  MAD_SD
## (Intercept) -0.1124  0.5214
## sexCode      0.7928  0.7001
##
## -----

```

```
## * For help interpreting the printed output see ?print.stanreg
## * For info on the priors used see ?prior_summary.stanreg
```

```
posterior_interval(logistic_sex)
```

```
##              5%      95%
## (Intercept) -0.9768360 0.7567365
## sexCode      -0.3335057 1.9450223
```

```
# As we can see, the log-odds coefficients, from which the probabilities derive, are
# insignificant as 0 is a plausible value here.
```

```
# Visually we can this here:
```

```
sims = as.matrix(logistic_sex)
```

```
int = sims[, 1]
```


```
female = int + sims[, 2]
```

```
plot(density(female), col = 'purple', lty = 2, lwd = 2,
     xlab = "log-odds Coefficient of Effect",
     main = "Posterior Density Function: Logistic Sex Model",
     y = "Probability Density", ylim = c(0, 1.5))
```

```
lines(density(int), lwd = 1, col = 'darkgreen')
```

```
legend('topleft', legend = c("Female", "Male"), col = c("purple", "darkgreen"), lty = 2:1)
```

Posterior Density Function: Logistic Sex Model



```
# We can see that the posterior distributions have very wide tails, indicating that
# the credible intervals are very large for our estimates. We can't, therefore, be sure
# in our estimated log effect of sex on testing positive for COVID. This is, again, good
# news for the prospect of absent confounding.
```



```
# t6 ~ LRTI/URTI/Control (Disease Classification)
g$statCode = ifelse(g$status == 'Asymptomatic', 0, ifelse(g$status == 'LRTI', 1, 2))
logistic_test = stan_glm(data = g, t6 ~ statCode, family = binomial(link = 'logit'))
```

```
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 1.9e-05 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.19 seconds.
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 1: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 1: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 1: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 1: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 1: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 1: Iteration:  1001 / 2000 [ 50%] (Sampling)
## Chain 1: Iteration:  1200 / 2000 [ 60%] (Sampling)
## Chain 1: Iteration:  1400 / 2000 [ 70%] (Sampling)
## Chain 1: Iteration:  1600 / 2000 [ 80%] (Sampling)
## Chain 1: Iteration:  1800 / 2000 [ 90%] (Sampling)
## Chain 1: Iteration:  2000 / 2000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.054286 seconds (Warm-up)
## Chain 1:                0.057757 seconds (Sampling)
## Chain 1:                0.112043 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 1.6e-05 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.16 seconds.
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 2: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 2: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 2: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 2: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 2: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 2: Iteration:  1001 / 2000 [ 50%] (Sampling)
## Chain 2: Iteration:  1200 / 2000 [ 60%] (Sampling)
## Chain 2: Iteration:  1400 / 2000 [ 70%] (Sampling)
## Chain 2: Iteration:  1600 / 2000 [ 80%] (Sampling)
## Chain 2: Iteration:  1800 / 2000 [ 90%] (Sampling)
## Chain 2: Iteration:  2000 / 2000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 0.054774 seconds (Warm-up)
## Chain 2:                0.060115 seconds (Sampling)
## Chain 2:                0.114889 seconds (Total)
```

```

## Chain 2:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 3).
## Chain 3:
## Chain 3: Gradient evaluation took 1.4e-05 seconds
## Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.14 seconds.
## Chain 3: Adjust your expectations accordingly!
## Chain 3:
## Chain 3:
## Chain 3: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 3: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 3: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 3: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 3: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 3: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 3: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 3: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 3: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 3: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 3: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 3: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 3:
## Chain 3: Elapsed Time: 0.05649 seconds (Warm-up)
## Chain 3:                0.056794 seconds (Sampling)
## Chain 3:                0.113284 seconds (Total)
## Chain 3:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 4).
## Chain 4:
## Chain 4: Gradient evaluation took 2.1e-05 seconds
## Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.21 seconds.
## Chain 4: Adjust your expectations accordingly!
## Chain 4:
## Chain 4:
## Chain 4: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 4: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 4: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 4: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 4: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 4: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 4: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 4: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 4: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 4: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 4: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 4: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 4:
## Chain 4: Elapsed Time: 0.049898 seconds (Warm-up)
## Chain 4:                0.054794 seconds (Sampling)
## Chain 4:                0.104692 seconds (Total)
## Chain 4:

```

```
print(logistic_test, digits = 4)
```

```
## stan_glm
## family:      binomial [logit]
## formula:     t6 ~ statCode
## observations: 36
## predictors:  2
## -----
##              Median  MAD_SD
## (Intercept) -1.0795  0.7922
## statCode     0.9895  0.4948
## -----
## * For help interpreting the printed output see ?print.stanreg
## * For info on the priors used see ?prior_summary.stanreg
```

```
posterior_interval(logistic_test)
```

```
##              5%      95%
## (Intercept) -2.5239562 0.1615144
## statCode     0.2178896 1.8861495
```

Given that there is a 95% probability that the effect of testing positive for COVID-19 increases, depending on whether someone is of classification URTI / LRTI, I will continue this analysis after controlling for dose concentration value.

```
logistic_char = stan_glm(data = g, t6 ~ statCode + conc,
                          family = binomial(link = 'logit'))
```

```
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 2e-05 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.2 seconds.
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 1: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 1: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 1: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 1: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 1: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 1: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 1: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 1: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 1: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 1: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.066513 seconds (Warm-up)
## Chain 1:                0.068595 seconds (Sampling)
## Chain 1:                0.135108 seconds (Total)
## Chain 1:
##
```

```

## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 1.5e-05 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.15 seconds.
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 2: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 2: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 2: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 2: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 2: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 2: Iteration:  1001 / 2000 [ 50%] (Sampling)
## Chain 2: Iteration:  1200 / 2000 [ 60%] (Sampling)
## Chain 2: Iteration:  1400 / 2000 [ 70%] (Sampling)
## Chain 2: Iteration:  1600 / 2000 [ 80%] (Sampling)
## Chain 2: Iteration:  1800 / 2000 [ 90%] (Sampling)
## Chain 2: Iteration:  2000 / 2000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 0.072704 seconds (Warm-up)
## Chain 2:                    0.071046 seconds (Sampling)
## Chain 2:                    0.14375 seconds (Total)
## Chain 2:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 3).
## Chain 3:
## Chain 3: Gradient evaluation took 1.5e-05 seconds
## Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.15 seconds.
## Chain 3: Adjust your expectations accordingly!
## Chain 3:
## Chain 3:
## Chain 3: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 3: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 3: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 3: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 3: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 3: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 3: Iteration:  1001 / 2000 [ 50%] (Sampling)
## Chain 3: Iteration:  1200 / 2000 [ 60%] (Sampling)
## Chain 3: Iteration:  1400 / 2000 [ 70%] (Sampling)
## Chain 3: Iteration:  1600 / 2000 [ 80%] (Sampling)
## Chain 3: Iteration:  1800 / 2000 [ 90%] (Sampling)
## Chain 3: Iteration:  2000 / 2000 [100%] (Sampling)
## Chain 3:
## Chain 3: Elapsed Time: 0.069716 seconds (Warm-up)
## Chain 3:                    0.065555 seconds (Sampling)
## Chain 3:                    0.135271 seconds (Total)
## Chain 3:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 4).
## Chain 4:
## Chain 4: Gradient evaluation took 1.5e-05 seconds
## Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.15 seconds.

```

```
## Chain 4: Adjust your expectations accordingly!
## Chain 4:
## Chain 4:
## Chain 4: Iteration: 1 / 2000 [ 0%] (Warmup)
## Chain 4: Iteration: 200 / 2000 [ 10%] (Warmup)
## Chain 4: Iteration: 400 / 2000 [ 20%] (Warmup)
## Chain 4: Iteration: 600 / 2000 [ 30%] (Warmup)
## Chain 4: Iteration: 800 / 2000 [ 40%] (Warmup)
## Chain 4: Iteration: 1000 / 2000 [ 50%] (Warmup)
## Chain 4: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 4: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 4: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 4: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 4: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 4: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 4:
## Chain 4: Elapsed Time: 0.064802 seconds (Warm-up)
## Chain 4: 0.073831 seconds (Sampling)
## Chain 4: 0.138633 seconds (Total)
## Chain 4:
```

```
print(logistic_char, digits = 4)
```

```
## stan_glm
## family: binomial [logit]
## formula: t6 ~ statCode + conc
## observations: 36
## predictors: 3
## -----
##           Median MAD_SD
## (Intercept) -0.0266 0.8798
## statCode    1.5708 0.6416
## conc        -5.9226 2.0820
##
## -----
## * For help interpreting the printed output see ?print.stanreg
## * For info on the priors used see ?prior_summary.stanreg
```

```
posterior_interval(logistic_char)
```

```
##           5%      95%
## (Intercept) -1.5321288 1.428489
## statCode    0.5758119 2.761981
## conc        -10.0377189 -2.960194
```

Interestingly, we still observe that there is a 95% probability that 0 is not a plausible value for the effect of disease classification on testing positive for COVID-19.

```
inv_logit(logistic_char$coefficients)
```

```
## (Intercept) statCode conc
## 0.493362024 0.827895800 0.002671007
```

```
# We can see that if someone is asymptomatic, the probability that they test positive is
# 0.496434705. This estimate is not significant, but our beta 1 value is. If someone suffers
# from LRTI, their probability of testing positive for COVID is predicted to be roughly
# 0.825825933.
```

```
inv_logit(2*logistic_char$coefficients[2])
```

```
## statCode
## 0.9585754
```

```
# Finally, if someone suffers from URTI, their probability of testing positive for COVID
# in this experiment will be the highest at 0.9574119.
```

```
# It is interesting to see that if you are an URTI patient, your probability of testing
# positive will be greater than if you were an LRTI patient, after controlling for the dose
# concentration differences the two between subjects.
```

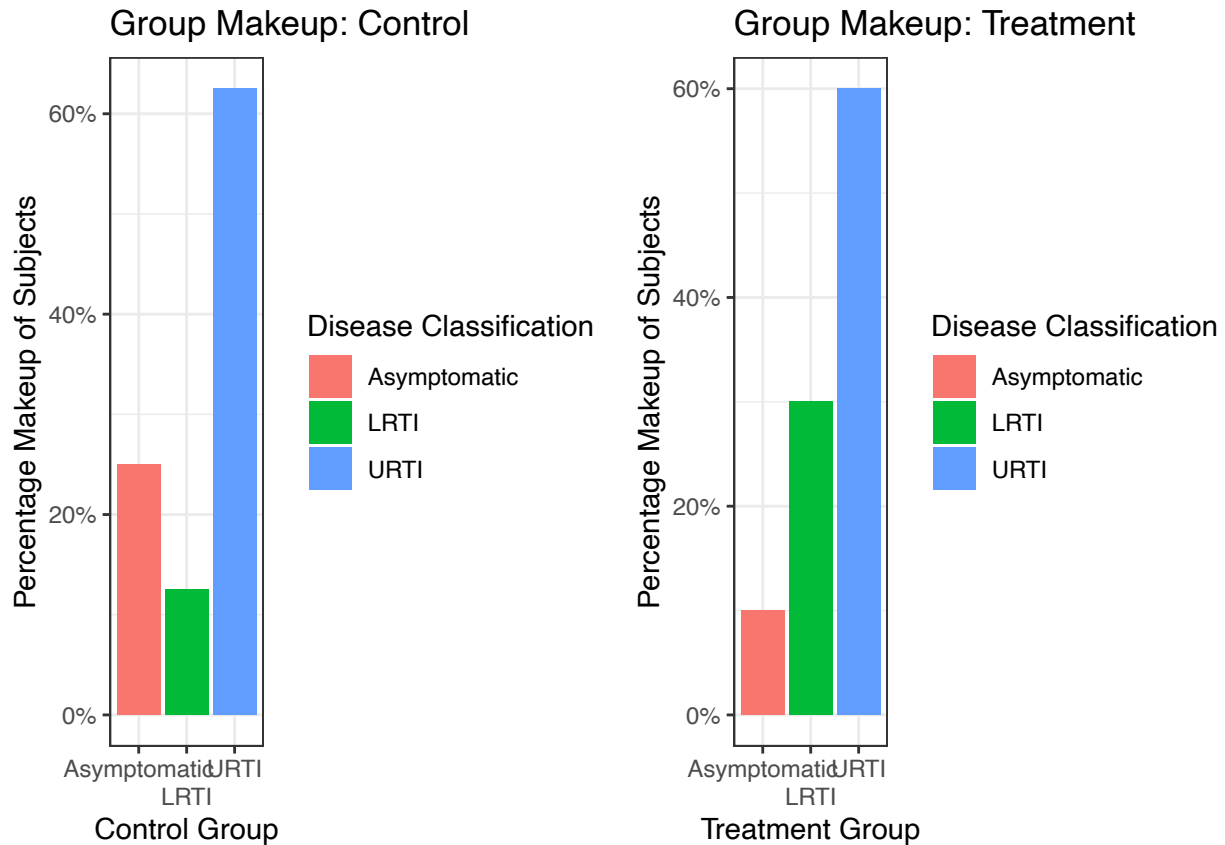
```
# This provides some interesting insight to the question of whether
# hydroxychloroquine & azithromycin are more effective towards LRTI patients.
```

```
# To check one more time, and determine whether there was a significant difference
# between the number of LTRI and HTRI patients in the treatment and control groups,
# we can visualize a bar chart of the counts.
```

```
g$treatment = ifelse(g$conc > 0, 'Treatment', 'Control')
df = tibble(g %>% group_by(treatment, status) %>%
  summarise(len = length(status)))
```

```
data1 = df %>% filter(treatment == 'Control') %>% mutate(prop = len/16)
data2 = df %>% filter(treatment == 'Treatment') %>% mutate(prop = len/20)
```

```
library(cowplot)
g1 = ggplot(data = data1, aes(x = status, y = prop, fill = as.factor(status))) +
  geom_col() + scale_y_continuous(labels = percent) +
  labs(x = 'Control Group',
    y = "Percentage Makeup of Subjects",
    title = "Group Makeup: Control") +
  scale_fill_discrete(name = 'Disease Classification') +
  scale_x_discrete(guide = guide_axis(n.dodge = 2))
g2 = ggplot(data = data2, aes(x = status, y = prop, fill = as.factor(status))) +
  geom_col() + scale_y_continuous(labels = percent) +
  scale_fill_discrete(name = 'Disease Classification') +
  labs(x = 'Treatment Group',
    y = "Percentage Makeup of Subjects",
    title = "Group Makeup: Treatment") +
  scale_x_discrete(guide = guide_axis(n.dodge = 2))
plot_grid(g1, g2)
```




There doesn't seem to be significantly more URTI/LRTI patients in the Treatment group than the Control group, scaled by their respective populations, however, this is an interesting question to ponder: whether hydroxychloroquine affects URTI and LRTI patients differently.

*# Finally, let's create a posterior density function of our parameter estimates: # * Remember this is after controlling for concentration effect.*

```
sims = as.matrix(logistic_char)
control = sims[, 1]
LRTI = sims[, 2] + control
URTI = 2*sims[, 2] + control
plot(density(LRTI), col = 'purple', lty = 2, lwd = 2,
     xlab = "Predicted log-odds Value of Testing Positive",
     main = "Posterior Density Function: Predicted Values",
     y = "Probability Density", ylim = c(0, 1), xlim = c(6, -3))

lines(density(control), lwd = 1, col = 'darkgreen')
lines(density(URTI), lwd = 2, col = 'darkorange', lty = 2)
legend('topleft', legend = c("LRTI Subjects", "Control", "URTI Subjects"),
     col = c("purple", "darkgreen", "darkorange"),
     lty = 2:1)
```

Posterior Density Function: Predicted Values



Given that 0 is not a major plausible value for the predicted log-odds value of LRTI & URTI patients, these predictions seem robust (but vary wildly in effect).

t6 ~ conc

`logistic_conc = stan_glm(data = g, t6 ~ conc, family = binomial(link = 'logit'))`

##

SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 2e-05 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.2 seconds.

Chain 1: Adjust your expectations accordingly!

Chain 1:

Chain 1:

Chain 1: Iteration: 1 / 2000 [0%] (Warmup)

Chain 1: Iteration: 200 / 2000 [10%] (Warmup)

Chain 1: Iteration: 400 / 2000 [20%] (Warmup)

Chain 1: Iteration: 600 / 2000 [30%] (Warmup)

Chain 1: Iteration: 800 / 2000 [40%] (Warmup)

Chain 1: Iteration: 1000 / 2000 [50%] (Warmup)

Chain 1: Iteration: 1001 / 2000 [50%] (Sampling)

Chain 1: Iteration: 1200 / 2000 [60%] (Sampling)

Chain 1: Iteration: 1400 / 2000 [70%] (Sampling)

Chain 1: Iteration: 1600 / 2000 [80%] (Sampling)

Chain 1: Iteration: 1800 / 2000 [90%] (Sampling)

Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)

Chain 1:

Chain 1: Elapsed Time: 0.054692 seconds (Warm-up)


```

## Chain 1:          0.059875 seconds (Sampling)
## Chain 1:          0.114567 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 1.8e-05 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.18 seconds.
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 2: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 2: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 2: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 2: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 2: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 2: Iteration:  1001 / 2000 [ 50%] (Sampling)
## Chain 2: Iteration:  1200 / 2000 [ 60%] (Sampling)
## Chain 2: Iteration:  1400 / 2000 [ 70%] (Sampling)
## Chain 2: Iteration:  1600 / 2000 [ 80%] (Sampling)
## Chain 2: Iteration:  1800 / 2000 [ 90%] (Sampling)
## Chain 2: Iteration:  2000 / 2000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 0.053842 seconds (Warm-up)
## Chain 2:          0.058586 seconds (Sampling)
## Chain 2:          0.112428 seconds (Total)
## Chain 2:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 3).
## Chain 3:
## Chain 3: Gradient evaluation took 3.1e-05 seconds
## Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.31 seconds.
## Chain 3: Adjust your expectations accordingly!
## Chain 3:
## Chain 3:
## Chain 3: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 3: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 3: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 3: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 3: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 3: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 3: Iteration:  1001 / 2000 [ 50%] (Sampling)
## Chain 3: Iteration:  1200 / 2000 [ 60%] (Sampling)
## Chain 3: Iteration:  1400 / 2000 [ 70%] (Sampling)
## Chain 3: Iteration:  1600 / 2000 [ 80%] (Sampling)
## Chain 3: Iteration:  1800 / 2000 [ 90%] (Sampling)
## Chain 3: Iteration:  2000 / 2000 [100%] (Sampling)
## Chain 3:
## Chain 3: Elapsed Time: 0.05562 seconds (Warm-up)
## Chain 3:          0.055722 seconds (Sampling)
## Chain 3:          0.111342 seconds (Total)
## Chain 3:
##

```

```
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 4).
## Chain 4:
## Chain 4: Gradient evaluation took 2.1e-05 seconds
## Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.21 seconds.
## Chain 4: Adjust your expectations accordingly!
## Chain 4:
## Chain 4:
## Chain 4: Iteration:    1 / 2000 [  0%] (Warmup)
## Chain 4: Iteration:   200 / 2000 [ 10%] (Warmup)
## Chain 4: Iteration:   400 / 2000 [ 20%] (Warmup)
## Chain 4: Iteration:   600 / 2000 [ 30%] (Warmup)
## Chain 4: Iteration:   800 / 2000 [ 40%] (Warmup)
## Chain 4: Iteration:  1000 / 2000 [ 50%] (Warmup)
## Chain 4: Iteration: 1001 / 2000 [ 50%] (Sampling)
## Chain 4: Iteration: 1200 / 2000 [ 60%] (Sampling)
## Chain 4: Iteration: 1400 / 2000 [ 70%] (Sampling)
## Chain 4: Iteration: 1600 / 2000 [ 80%] (Sampling)
## Chain 4: Iteration: 1800 / 2000 [ 90%] (Sampling)
## Chain 4: Iteration: 2000 / 2000 [100%] (Sampling)
## Chain 4:
## Chain 4: Elapsed Time: 0.055226 seconds (Warm-up)
## Chain 4:                0.059854 seconds (Sampling)
## Chain 4:                0.11508 seconds (Total)
## Chain 4:
```

```
print(logistic_conc, digits = 4)
```

```
## stan_glm
## family:      binomial [logit]
## formula:     t6 ~ conc
## observations: 36
## predictors:  2
## -----
##              Median MAD_SD
## (Intercept)  1.6195  0.6180
## conc         -4.2418  1.5470
## -----
## * For help interpreting the printed output see ?print.stanreg
## * For info on the priors used see ?prior_summary.stanreg
```

```
posterior_interval(logistic_conc)
```

```
##              5%          95%
## (Intercept)  0.6768372  2.788618
## conc         -7.1112006 -1.838666
```

```
# There is a 95% probability that the true log-odds effect of hydroxychloroquine
# concentration on testing positive for COVID-19 is between -6.9077346 & -1.939355
# *This is a significant finding, and is far larger than anything we have observed from
# the other variables.
```

```
inv_logit(logistic_conc$coefficients[1])
```

```
## (Intercept)  
## 0.8347258
```

```
# For a probabilistic interpretation: someone without the dose concentration has a 0.8333466  
# probability of testing positive for COVID.
```

```
inv_logit(logistic_conc$coefficients[1] + .5*logistic_conc$coefficients[2])
```

```
## (Intercept)  
## 0.3772145
```

```
# If someone increases their dose by .5 ( $\mu$  g/ml), their probability of testing positive is  
# 0.3789761. As we can see, this is significantly lower.
```

```
# Onto an additional method:
```

```
# Although this analysis has rested on Bayesian methods, I feel that it's useful, given the  
# relatively small sample size of subjects in this clinical experiment, to conduct frequentist  
# bootstrap analysis.
```

```
# Using this method, we can sample from our sample (with replacement) using our original  
# sample data. Given that we assume our sample is our best estimate of what our population  
# distribution looks like, sampling from it is akin to running a continuous amount of  
# homogenous clinical trials, all outcomes in some way representing what our population  
# results would be.
```

```
# It is important to note that this bootstrap method outlined is based on the Frequentist  
# school of thought. I've decided to briefly switch over from Bayesian methods, as I don't  
# have any real experience with bayesian bootstrapping, and I feel as if I wouldn't be  
# accurately able to explain what is going on behind the code. My understanding is our  
# current described method wouldn't make intuitive sense to a Bayesian, as the data is  
# presumed fixed, and repeatedly sampling from it would acknowledge that it comes from a  
# larger population distribution and is not "fixed." In this sense, a bayesian bootstrap  
# would strive to sample from a series of posterior distributions containing regression  
# coefficients, instead of the observed data sample.
```

```
set.seed(101)  
library(mosaic)  
lower = c()  
higher = c()  
coef = c()  
sims = 1000
```

```
for (i in 1:sims) {  
  new = resample(g)  
  m = glm(data = new, t6 ~ conc, family = binomial(link = 'logit'))  
}
```


```

coef[i] = m$coefficients[2]
lower[i] = suppressMessages(confint(m))[2]
higher[i] = suppressMessages(confint(m))[4]
}

hist(coef, 30, col = 'darkorange', prob = T,
     xlab = "Dose Concentration Effect on Testing Positive",
     main = 'Histogram of Dose Concentration ( $\mu$  g/ml) Effect')
curve(dnorm(x, mean = mean(coef), sd = sd(coef)), add = T, col = 'purple', lwd = 2)
abline(v = mean(coef), col = 'darkblue', lwd = 2)
abline(v = mean(lower), col = 'darkgreen', lwd = 2, lty = 2)
abline(v = mean(higher), col = 'darkgreen', lwd = 2, lty = 2)
legend('topleft', legend = c("Mean Effect", "95% Confidence Interval"),
     col = c("darkblue", "darkgreen"), lty = 1:2, cex = .75)

```

Histogram of Dose Concentration (μ g/ml) Effect



```

bootstrap_CI = c(mean(lower), mean(higher))
bootstrap_CI

```

```
## [1] -8.312200 -1.522801
```

*# 95% Frequentist Confidence Interval from bootstrap sampling:
 # We are 95% certain that the true, log-odds effect of hydroxychloroquine concentration on
 # testing positive for COVID lies between (-8.312200, -1.522801).*

*# Rather than being based on the observed sample data, which is assumed to be "fixed" from
 # our above observations, this result is based on hypothesized repeated experiments.
 # This is an important distinction, as we assume that our data is 'random' in that our*

```
# observations are sampled from a larger, unknown population distribution. We then try to
# estimate a point estimate from modeling our regression estimators after our sample data, in
# contrast to the bayesian assumption that our parameter estimates follow a distribution and
# we can assume a prior distribution before attempting to maximize the likelihood of observing
# the data given our model parameters. This is the reason why we are allowed to estimate a
# posterior distribution (a compromise between our prior distribution assumptions and the
# evidence of what our effect looks like after having seen our fixed data sample) of the
# coefficients above.
```

```
# In this frequentist interpretation, we can say that if we were to repeat this experiment
# an infinite amount of times and draw confidence intervals each of the time for dose effect,
# 95% of those intervals will contain the true, parameter effect of dose concentration value
# on testing positive for COVID.
```

```
# * Remember, these effects resemble log-odds, as we can't gauge whether 0 is a plausible
# value if we convert the log-odds values into probabilities. This is because the interval
# can never encompass 0 with both negative and positive values, as non-negativity is a
# self-evident truth of probability, written into Kolmogorov's first axiom that  $P(A) \geq 0$ .
```

```
# While we are using a Frequentist method, let's analyze statistical significance through the
# p-value:
```

```
z = (0 - mean(coef)) / sd(coef)
p = pnorm(z, lower.tail = T)
qnorm(p)
```

```
## [1] 2.473545
```

```
sum(coef >= 0) / sims
```

```
## [1] 0.001
```

```
# This is our p-value, which tells us the probability that we have observed what we have
# observed, given that our null hypothesis is true. More specifically, it is the probability
# that we observe an equal or greater value than our test-statistic of roughly ~ 2.5,
# assuming that the beta effect is 0.
```

```
# As we can see, the probability of observing this value, under the assumption that our
# null hypothesis is true, is extremely low. We can, thus, reject the null hypothesis of a 0
# beta value, for the dose concentration, log effect of testing positive for COVID,
# at the 5% & 1% significance level.
```

```
# For a visual representation of the logistic regression:
```

```
sims = as.matrix(logistic_conc)
int = sims[, 1]
conc = sims[, 2]
```

```
lower = c()
```


```

higher = c()
med = c()
sequence = seq(min(g$conc)-.5, max(g$conc)+.5, .001)

for (i in 1:length(sequence)) {
  ypred = int + conc*sequence[i]
  lower[i] = quantile(ypred, prob = .025)
  higher[i] = quantile(ypred, prob = .975)
  med[i] = median(ypred)
}

plot(g$conc, g$t6, pch = 16, xlab = "Concentration of Hydroxychloroquine ( $\mu$  g/ml)",
     ylab = "Probability of Testing Positive for COVID-19",
     main = "Logistic Regression of COVID Testing Outcome on Dose Concentration")
lines(sequence, inv_logit(lower), lty = 2, col = "darkblue")
lines(sequence, inv_logit(higher), lty = 2, col = "darkblue")
lines(sequence, inv_logit(med), lwd = 2, col = "magenta")

```



*# It is nice to see that our logistic regression now follows a sigmoid curve, thus indicating
that hydroxychloroquine dose concentration serves as a robust probabilistic indicator for
COVID-19 test results*

*# This chunk of code will be attributed to result recreation in the original
experiment.*

```

library(readxl)
library(sqldf)
virus = read_excel("/Users/patrickpoleshuk/gautretData_(1).xlsx")

```

```

virus$yes = ifelse(virus$chloroquine == "Yes", 1, 0)

outlist = c()
fn <- function(x){
  outlist <- c(outlist, mean(x), sd(x))
}
results = tapply(virus$age, virus$yes, fn)
#results_2 = as.data.frame(tapply(virus$age, virus$yes, fn))
results = as.data.frame(do.call(rbind, results))
results[1] = formatC(results$V1, digits = 2)
results[2] = formatC(results$V2, digits = 2)

names(results) <- c("Age_Mean", "Age_SD")
info = sqldf("SELECT Age_Mean || ' +/- ' || Age_SD AS concat FROM results")
de <- as.data.frame(sqldf("SELECT ROUND(avg(age),2) || ' +/- ' ||
                          ROUND(stdev(age), 2) AS concat FROM virus"))
info = rbind(info, de)

with(virus, t.test(age[yes == 0], age[yes == 1]))

```

```

##
## Welch Two Sample t-test
##
## data: age[yes == 0] and age[yes == 1]
## t = -1.893, df = 27.9, p-value = 0.06878
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -28.787767 1.137767
## sample estimates:
## mean of x mean of y
## 37.375 51.200

```

```

# p-value = 0.06878
# t-value = -1.893
info = cbind(info, "-1.9", ".06")
info[1, 2] = ""; info[3, 2] = ""
info[1, 3] = ""; info[3, 3] = ""

extra = as.data.frame(as.matrix(c("Control(N=16)",
"Hydroxychloroquine(N=20)", "All(36)")))
info = cbind(extra, info)

gen = sqldf("SELECT COUNT(sex) FROM virus WHERE sex = 'M'
            GROUP BY chloroquine")

library(qpcR)
info = qpcR:::cbind.na(info, gen)
info[3, NCOL(info)] <- sum(gen$`COUNT(sex)`)
with(virus, t.test(sexCode[yes == 0], sexCode[yes == 1]))

```

```

##
## Welch Two Sample t-test

```

```
##
## data:  sexCode[yes == 0] and sexCode[yes == 1]
## t = 0.44309, df = 32.566, p-value = 0.6606
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -0.2695505  0.4195505
## sample estimates:
## mean of x mean of y
##      0.625      0.550
```

```
# p-value = 0.6606
```

```
info = cbind(info, ".66")
info[1, 6] = ""; info[3, 6] = ""

clin_1 = sqldf("
  SELECT COUNT(status) FROM virus
  WHERE status = 'Asymptomatic'
  GROUP BY chloroquine
")

clin_2 = sqldf("
  SELECT COUNT(status) st2 FROM virus
  WHERE status = 'URTI'
  GROUP BY chloroquine
")

clin_3 = sqldf("
  SELECT COUNT(status) st3 FROM virus
  WHERE status = 'LRTI'
  GROUP BY chloroquine
")

info = qpcR::cbind.na(info, clin_1, clin_2, clin_3)
info[3, 7] = sum(clin_1$`COUNT(status)`)
info[3, 8] = sum(clin_2$st2)
info[3, 9] = sum(clin_3$st3)

v = sqldf("SELECT * FROM virus WHERE status IS NOT NULL")
v$status = as.factor(v$status)
v$status = as.numeric(v$status)
s = v %>% filter(status != 1)
with(s, t.test(status[yes == 0], status[yes == 1]))
```

```
##
## Welch Two Sample t-test
##
## data:  status[yes == 0] and status[yes == 1]
## t = 1.0397, df = 26.906, p-value = 0.3077
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -0.1623087  0.4956420
## sample estimates:
## mean of x mean of y
```



```
## 2.833333 2.666667
```

```
# p-value = 0.3077  
# We found that this p-value is calculated by filtering out the  
# asymptomatic classifications, and comparing if LRTI & URTI classes  
# are disproportionately assigned to the treatment/control group.  
# We see, overall, from the table that this is not true in any of the  
# cases, except for age which the study has already established is  
# different, with the treatment group being older on average.
```

```
info = cbind(info, ".31")  
info[1, 10] = ""; info[3, 10] = ""
```

```
names(virus)[4] <- c("inclusion_time")
```

```
inc = sqldf("SELECT ROUND(avg(inclusion_time), 2),  
ROUND(stdev(inclusion_time), 2) FROM virus  
GROUP BY chloroquine ")
```

```
v = sqldf("SELECT inclusion_time, yes FROM virus WHERE  
inclusion_time IS NOT NULL AND yes IS NOT NULL")  
v$inclusion_time = as.numeric(v$inclusion_time)
```

```
ans = sqldf("SELECT ROUND(avg(inclusion_time), 2), ROUND(stdev(inclusion_time), 2) FROM v  
GROUP BY yes")
```

```
with(v, t.test(inclusion_time[yes == 0], inclusion_time[yes == 1]))
```

```
##  
## Welch Two Sample t-test  
##  
## data: inclusion_time[yes == 0] and inclusion_time[yes == 1]  
## t = -0.14381, df = 17.651, p-value = 0.8873  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
## -2.431224 2.120113  
## sample estimates:  
## mean of x mean of y  
## 3.900000 4.055556
```

```
# t-value = -0.14381  
# p-value = 0.8873  
info = qpcR::cbind.na(info, ans)  
  
names(info)[11:12] <- c("mi", "si")  
v = sqldf("SELECT * FROM v WHERE inclusion_time IS NOT NULL")  
info[3, 11] <- formatC(mean(v$inclusion_time), digits = 2)  
info[3, 12] <- formatC(sd(v$inclusion_time), digits = 2)  
  
concat2 = sqldf("SELECT mi || ' +/- ' || si AS concat2 FROM info")  
info = cbind(info, concat2)  
info = info %>% dplyr::select(-mi, -si)
```

```

info = cbind(info, "-.14", ".88")
info[1, 12] = ""; info[3, 12] = ""
info[1, 13] = ""; info[3, 13] = ""

names(info) <- c("", "M+-SD[AGE]", "t",
                "p", "Male-count", "p",
                "Asymp", "URTI", "LRTI",
                "p", "M+-S[INCLU]", "t",
                "p")

library(xtable)
print.xtable(xtable(info), file = "./Downloads/info_covid_2.txt")

agrestiCaffo <- function(x1,n1,x2,n2,alpha=.05) {
  #
  # The Agresti-Caffo proportion estimate with alpha = .05
  #

  zCrit      <- qnorm(1-alpha/2);

  pHat1      <- (x1)/(n1);
  pPrime1    <- (x1+1)/(n1+2);
  pPrimeV1   <- (pPrime1*(1-pPrime1)/(n1+2));
  pHat2      <- (x2)/(n2);
  pPrime2    <- (x2+1)/(n2+2);
  pPrimeV2   <- (pPrime2*(1-pPrime2)/(n2+2));
  pPrimeSD   <- sqrt(pPrimeV1+pPrimeV2);
  lCI        <- pPrime1-pPrime2 - zCrit*pPrimeSD;
  uCI        <- pPrime1-pPrime2 + zCrit*pPrimeSD;
  zScore     <- (pPrime1-pPrime2)/pPrimeSD;

  pValue     <- (1-pnorm(abs(zScore)))*2;
  pValueGgtP <- (1-pnorm(zScore))

  myParms    <- list()

  myParms$zCrit    <- zCrit;
  myParms$pHat1    <- pHat1;
  myParms$pPrime1  <- pPrime1;
  myParms$pHat2    <- pHat2;
  myParms$pPrime2  <- pPrime2;
  myParms$deltaP   <- pHat1-pHat2;
  myParms$lower    <- lCI;
  myParms$upper    <- uCI;
  myParms$pValue   <- pValue;
  myParms$pValueGgtP <- pValueGgtP;
  return(myParms)
}

# DAY 3:
sqldf("SELECT COUNT(D3) FROM virus WHERE chloroquine = 'Yes' AND
      D3 = 'NEG' ")

```

```
##      COUNT(D3)
## 1          10
```

```
# 10
sqldf("SELECT COUNT(D3) FROM virus WHERE chloroquine = 'Yes'")
```

```
##      COUNT(D3)
## 1          20
```

```
# 20
sqldf("SELECT COUNT(D3) FROM virus WHERE chloroquine = 'No' AND
      D3 = 'NEG' ")
```

```
##      COUNT(D3)
## 1           1
```

```
# 1
sqldf("SELECT COUNT(D3) FROM virus WHERE chloroquine = 'No'")
```

```
##      COUNT(D3)
## 1          16
```

```
# 16
# DAY 4:
sqldf("SELECT COUNT(D4) FROM virus WHERE chloroquine = 'Yes' AND
      D4 = 'NEG' ")
```

```
##      COUNT(D4)
## 1          12
```

```
# 12
sqldf("SELECT COUNT(D4) FROM virus WHERE chloroquine = 'Yes'")
```

```
##      COUNT(D4)
## 1          20
```

```
# 20
sqldf("SELECT COUNT(D4) FROM virus WHERE chloroquine = 'No' AND
      D4 = 'NEG'")
```

```
##      COUNT(D4)
## 1           4
```

```
# 4
sqldf("SELECT COUNT(D4) FROM virus WHERE chloroquine = 'No'")
```

```
##      COUNT(D4)
## 1          16
```

```
# 16
```

```
# NOTE: We count ND for the negative treatment count on day 5 & day 6,  
# because even though the subject hasn't been tested that day, it was already  
# shown in the previous day that the subject tested negative.
```

```
# DAY 5:
```

```
sqldf("SELECT COUNT(D5) FROM virus WHERE chloroquine = 'Yes' AND  
      (D5 = 'ND' OR D5 = 'NEG')")
```

```
##      COUNT(D5)  
## 1          13
```

```
# 13
```

```
sqldf("SELECT COUNT(D5) FROM virus WHERE chloroquine = 'Yes'")
```

```
##      COUNT(D5)  
## 1          20
```

```
# 20
```

```
sqldf("SELECT COUNT(D5) FROM virus WHERE chloroquine = 'No' AND  
      D5 = 'NEG'")
```

```
##      COUNT(D5)  
## 1           3
```

```
# 3
```

```
sqldf("SELECT COUNT(D5) FROM virus WHERE chloroquine = 'No'")
```

```
##      COUNT(D5)  
## 1          16
```

```
# 16
```

```
# DAY 6:
```

```
sqldf("SELECT COUNT(D6) FROM virus WHERE chloroquine = 'Yes' AND  
      (D6 = 'NEG' OR D6 = 'ND')")
```

```
##      COUNT(D6)  
## 1          14
```

```
# 14
```

```
sqldf("SELECT COUNT(D6) FROM virus WHERE chloroquine = 'Yes'")
```

```
##      COUNT(D6)  
## 1          20
```

```
# 20
```

```
sqldf("SELECT COUNT(D6) FROM virus WHERE chloroquine = 'No' AND  
      D6 = 'NEG'")
```

```
##      COUNT(D6)  
## 1          2
```

```
# 2
```

```
sqldf("SELECT COUNT(D6) FROM virus WHERE chloroquine = 'No'")
```

```
##      COUNT(D6)  
## 1          16
```

```
# 16
```

```
agrestiCaffo(10, 20, 1, 16)
```

```
## $zCrit  
## [1] 1.959964  
##  
## $pHat1  
## [1] 0.5  
##  
## $pPrime1  
## [1] 0.5  
##  
## $pHat2  
## [1] 0.0625  
##  
## $pPrime2  
## [1] 0.1111111  
##  
## $deltaP  
## [1] 0.4375  
##  
## $lower  
## [1] 0.1344662  
##  
## $upper  
## [1] 0.6433116  
##  
## $pValue  
## [1] 0.002736951  
##  
## $pValueGtP  
## [1] 0.001368476
```

```
# 0.002736951
```

```
agrestiCaffo(12, 20, 4, 16)
```

```
## $zCrit
## [1] 1.959964
##
## $pHat1
## [1] 0.6
##
## $pPrime1
## [1] 0.5909091
##
## $pHat2
## [1] 0.25
##
## $pPrime2
## [1] 0.2777778
##
## $deltaP
## [1] 0.35
##
## $lower
## [1] 0.02154174
##
## $upper
## [1] 0.6047209
##
## $pValue
## [1] 0.0353122
##
## $pValueGtP
## [1] 0.0176561
```

```
# 0.0353122
agrestiCaffo(13, 20, 3, 16)
```

```
## $zCrit
## [1] 1.959964
##
## $pHat1
## [1] 0.65
##
## $pPrime1
## [1] 0.6363636
##
## $pHat2
## [1] 0.1875
##
## $pPrime2
## [1] 0.2222222
##
## $deltaP
## [1] 0.4625
##
## $lower
## [1] 0.1361262
##
```

```
## $upper
## [1] 0.6921566
##
## $pValue
## [1] 0.003504445
##
## $pValueGgtP
## [1] 0.001752222
```

```
# 0.003504445
agrestiCaffo(14, 20, 2, 16)
```

```
## $zCrit
## [1] 1.959964
##
## $pHat1
## [1] 0.7
##
## $pPrime1
## [1] 0.6818182
##
## $pHat2
## [1] 0.125
##
## $pPrime2
## [1] 0.1666667
##
## $deltaP
## [1] 0.575
##
## $lower
## [1] 0.2553024
##
## $upper
## [1] 0.7750006
##
## $pValue
## [1] 0.0001020631
##
## $pValueGgtP
## [1] 5.103155e-05
```

```
# .0001020631
```

```
library(tidyverse)
table2 = tibble(
  x = c("Hydroxychloroquine(N=20)",
"Control(N=16)"), Prop_Day3 = c('10/20',
                                     '1/16'), Percent_Day3 =
  c('50%', '6.3%'), p3 = c("", ".003"), Prop_Day4 = c('12/20',
                                     '4/16'),
  Percent_Day4 = c('60%', '25%'), p4 = c("", "0.04"), Prop_Day5 =
  c('13/20', '3/16'), Percent_Day5 = c("56%", "18.8%"), p5 = c(
```

```

    "", ".004"), Prop_Day6 = c("14/20", "2/16"), Percent_Day6 = c(
      "70%", "12.5%"), p6 = c("", ".0001")
)

names(table2) <- c("", "Neg/Pos3", "%", "p3", "Neg/Pos4", "%", "p4",
  "Neg/Pos5", "%", "p5", "Neg/Pos6", "%", "p6")

print.xtable(xtable(table2), file = "./table2_covid_data.txt")

# With the Azithromycin counts, we don't need to do new counts for
# the control group and chloroquine tests. Since only the subjects who
# were dosed with Azithromycin were in the treatment group, the negative
# counts can just be subtracted from the base hydroxychloroquine negative
# counts.

# To find the statistical significance of the differences between
# three proportions, I just pooled the Azithromycin & Hydroxychloroquine
# proportions against the control proportions.

# DAY 3:
sqldf("SELECT COUNT(D3) FROM virus WHERE azithromycin = 'Yes' AND
      D3 = 'NEG' ")

```

```

##      COUNT(D3)
## 1              5

```

```

# 5
sqldf("SELECT COUNT(D3) FROM virus WHERE azithromycin = 'Yes'")

```

```

##      COUNT(D3)
## 1              6

```

```

# 6

# DAY 4:
sqldf("SELECT COUNT(D4) FROM virus WHERE azithromycin = 'Yes' AND
      D4 = 'NEG' ")

```

```

##      COUNT(D4)
## 1              5

```

```

# 5
sqldf("SELECT COUNT(D4) FROM virus WHERE azithromycin = 'Yes'")

```

```

##      COUNT(D4)
## 1              6

```

```

# 6

# DAY 5:
sqldf("SELECT COUNT(D5) FROM virus WHERE azithromycin = 'Yes' AND
      D5 = 'NEG' ")

```



```
## COUNT(D5)
## 1 6
```

```
# 6
sqldf("SELECT COUNT(D5) FROM virus WHERE azithromycin = 'Yes'")
```

```
## COUNT(D5)
## 1 6
```

```
# DAY 6:
sqldf("SELECT COUNT(D6) FROM virus WHERE azithromycin = 'Yes' AND
      D6 = 'NEG' ")
```

```
## COUNT(D6)
## 1 6
```

```
# 6
sqldf("SELECT COUNT(D6) FROM virus WHERE azithromycin = 'Yes'")
```

```
## COUNT(D6)
## 1 6
```

```
# 6

# Pooled  $5/14 + 5/6 = 10/20$ 
agrestiCaffo(1, 16, 10, 20)
```

```
## $zCrit
## [1] 1.959964
##
## $pHat1
## [1] 0.0625
##
## $pPrime1
## [1] 0.1111111
##
## $pHat2
## [1] 0.5
##
## $pPrime2
## [1] 0.5
##
## $deltaP
## [1] -0.4375
##
## $lower
## [1] -0.6433116
##
## $upper
## [1] -0.1344662
##
```

```
## $pValue
## [1] 0.002736951
##
## $pValueGgtP
## [1] 0.9986315
```

```
# 0.002736951
```

```
# Pooled  $7/14 + 5/6 = 12/20$ 
agrestiCaffo(4, 16, 12, 20)
```

```
## $zCrit
## [1] 1.959964
##
## $pHat1
## [1] 0.25
##
## $pPrime1
## [1] 0.2777778
##
## $pHat2
## [1] 0.6
##
## $pPrime2
## [1] 0.5909091
##
## $deltaP
## [1] -0.35
##
## $lower
## [1] -0.6047209
##
## $upper
## [1] -0.02154174
##
## $pValue
## [1] 0.0353122
##
## $pValueGgtP
## [1] 0.9823439
```

```
# 0.0353122
```

```
# Pooled  $7/14 + 6/6 = 13/20$ 
agrestiCaffo(3, 16, 13, 20)
```

```
## $zCrit
## [1] 1.959964
##
## $pHat1
## [1] 0.1875
##
## $pPrime1
```

```
## [1] 0.2222222
##
## $pHat2
## [1] 0.65
##
## $pPrime2
## [1] 0.6363636
##
## $deltaP
## [1] -0.4625
##
## $lower
## [1] -0.6921566
##
## $upper
## [1] -0.1361262
##
## $pValue
## [1] 0.003504445
##
## $pValueGtP
## [1] 0.9982478
```

```
# 0.003504445
```

```
# Pooled  $8/14 + 6/6 = 14/20$ 
agrestiCaffo(2, 16, 14, 20)
```

```
## $zCrit
## [1] 1.959964
##
## $pHat1
## [1] 0.125
##
## $pPrime1
## [1] 0.1666667
##
## $pHat2
## [1] 0.7
##
## $pPrime2
## [1] 0.6818182
##
## $deltaP
## [1] -0.575
##
## $lower
## [1] -0.7750006
##
## $upper
## [1] -0.2553024
##
## $pValue
## [1] 0.0001020631
```

```
##
## $pValueGgtP
## [1] 0.999949
```

```
# 0.0001020631
```

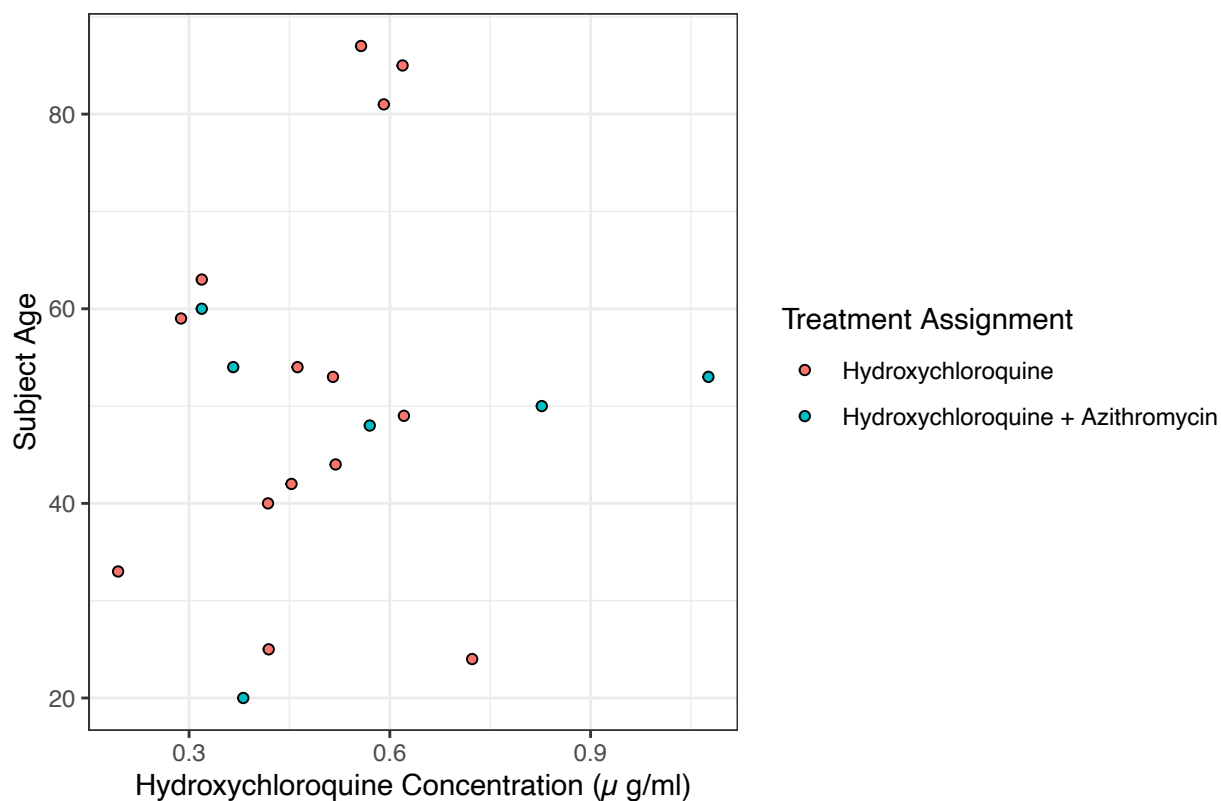
```
table3 = tibble(
  x = c("Control(N=16)",
    "HydC.(N=14)", "Azithromycin(N=6)"), Prop_Day3 = c('1/16',
    '5/14', "5/6"),
  Percent_Day3 =
    c('6.3%', '35.7%', '83.3%'), p3 = c("", ".003", ""),
  Prop_Day4 = c('4/16', '7/14', '5/6'),
  Percent_Day4 = c('25%', '50%', '83.3%'), p4 = c("", "0.04", ""),
  Prop_Day5 =
    c('3/16', '7/14', '6/6'), Percent_Day5 = c("18.8%", "50%", "100%"),
  p5 = c(
    "", ".004", ""), Prop_Day6 = c("2/16", "8/14", "6/6"),
  Percent_Day6 = c(
    "12.5%", "57.1%", "100%"), p6 = c("", ".0001", "")
)
```

```
names(table3) <- c("", "Neg/Pos3", "%", "p3", "Neg/Pos4", "%", "p4",
  "Neg/Pos5", "%", "p5", "Neg/Pos6", "%", "p6")
print.xtable(xtable(table3), file = "./table3_covid_data.txt")
```

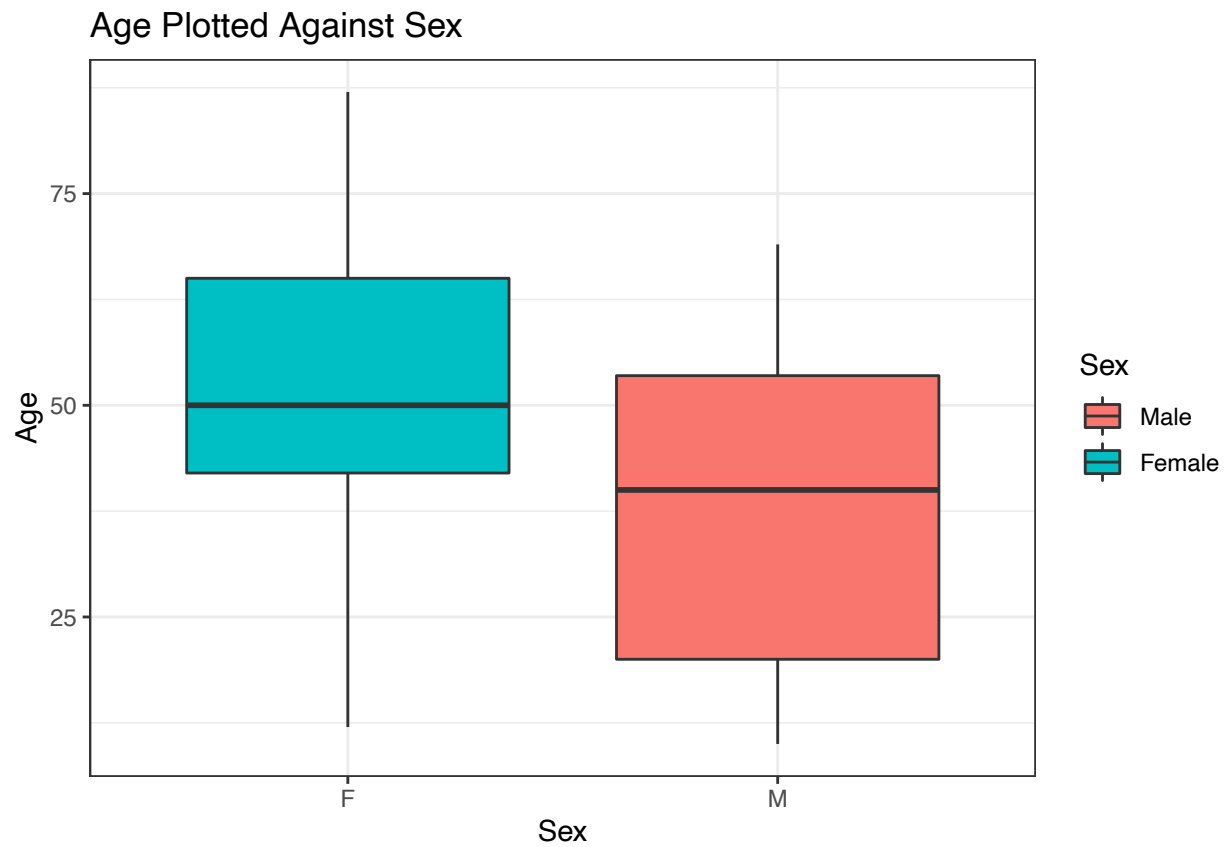
```
# Additional Graphics for Potential Variable Relationships.
```

```
gg = as.data.frame(g %>% filter(conc > 0))
gg$code = ifelse(gg$category == "Control", 0, ifelse(gg$category == "Hydrox", 1, 2))
ggplot(data = gg, aes(y = age, x = conc, fill = as.factor(as.numeric(code)))) +
  geom_point(pch = 21) +
  labs(title = "Age Plotted Against Hydroxychloroquine Concentration",
    x = "Hydroxychloroquine Concentration ( $\mu$  g/ml)",
    y = "Subject Age") + scale_fill_discrete(name = 'Treatment Assignment',
    labels = c("Hydroxychloroquine",
      "Hydroxychloroquine + Azithromycin"))
```

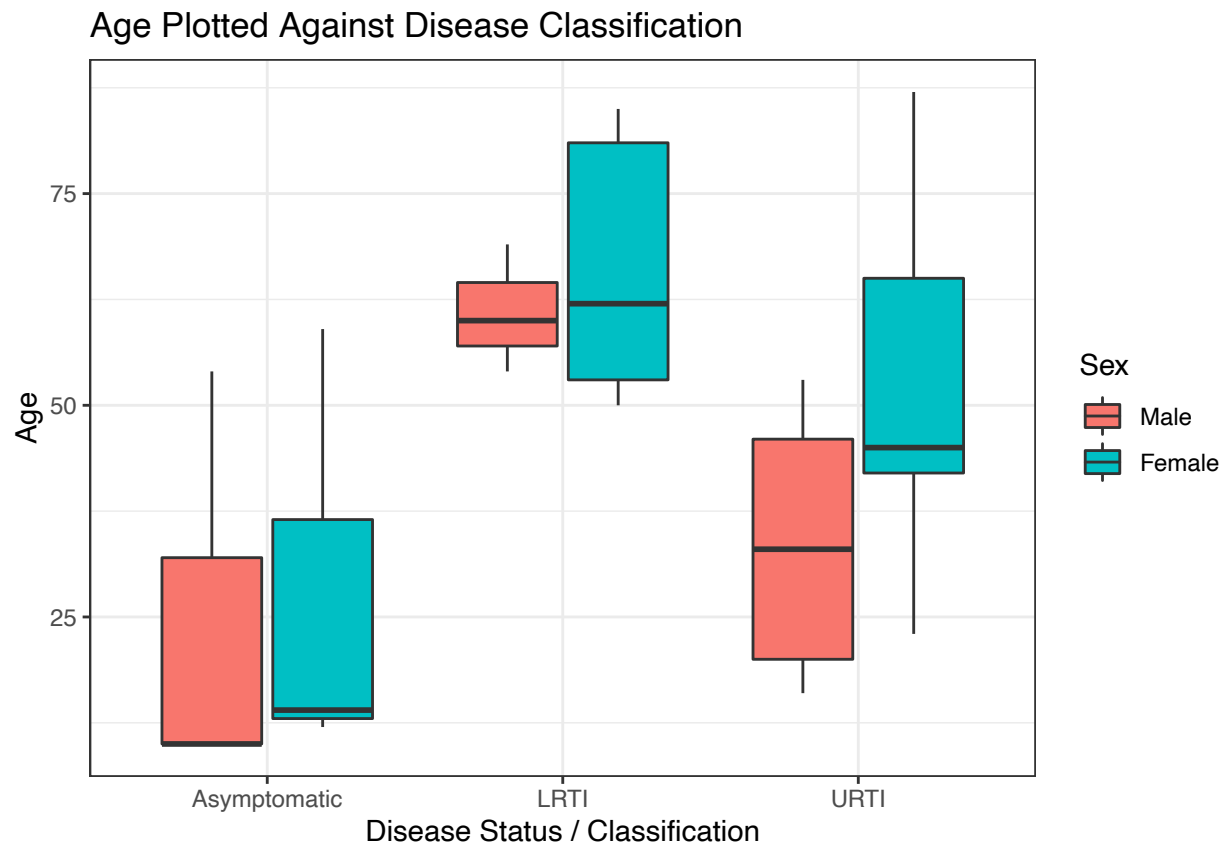
Age Plotted Against Hydroxychloroquine Concentration



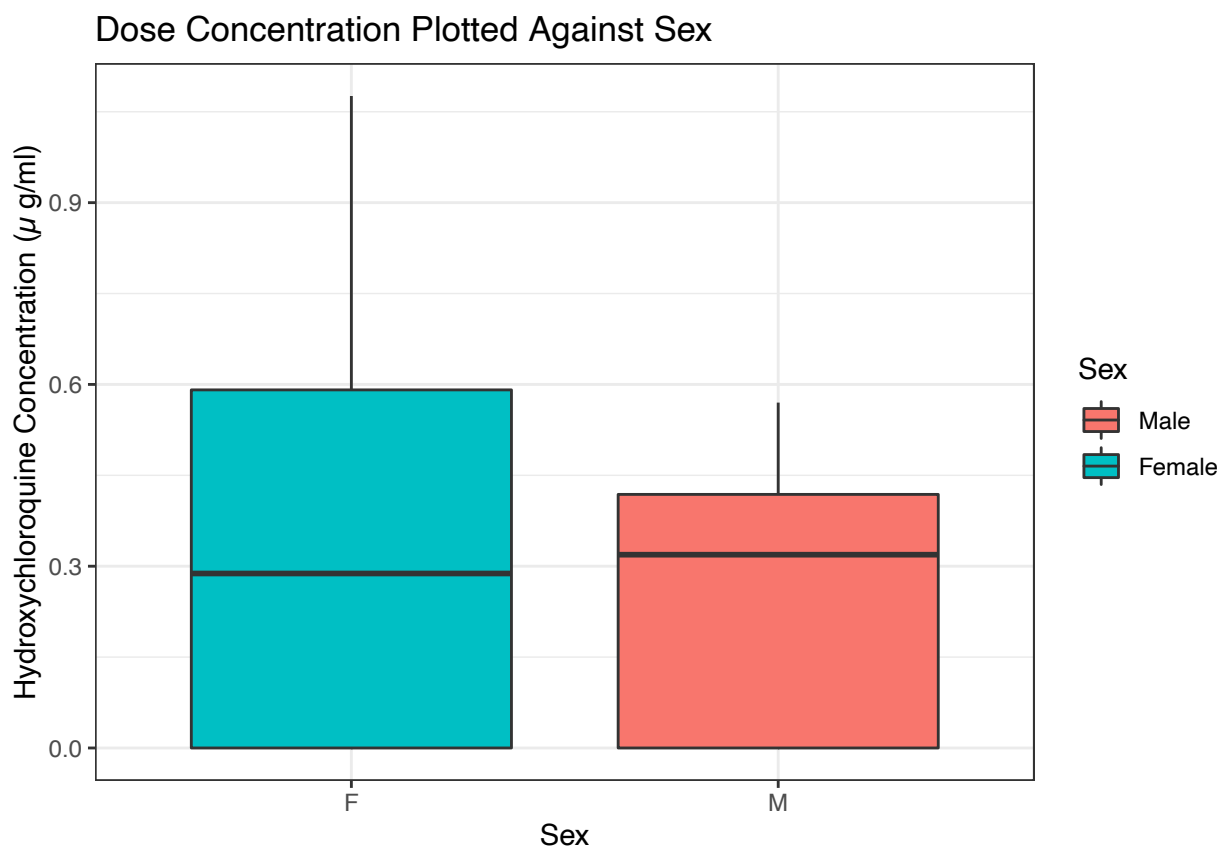
```
ggplot(data = g, aes(y = age, x = sex, fill = as.factor(sexCod))) +
  geom_boxplot() + labs(x = "Sex", y = "Age",
    title = "Age Plotted Against Sex") +
  scale_fill_discrete(name = 'Sex', labels = c("Male", "Female"))
```



```
ggplot(data = g, aes(y = age, x = status, fill = as.factor(sexCODE))) +  
  geom_boxplot() + labs(x = "Disease Status / Classification", y = "Age",  
                        title = "Age Plotted Against Disease Classification") +  
  scale_fill_discrete(name = 'Sex', labels = c('Male', "Female"))
```



```
ggplot(data = g, aes(y = conc, x = sex, fill = as.factor(sexCODE))) +
  geom_boxplot() + labs(x = "Sex",
                        y = "Hydroxychloroquine Concentration (µ g/ml)",
                        title = "Dose Concentration Plotted Against Sex") +
  scale_fill_discrete(name = 'Sex', labels = c('Male', 'Female'))
```



```
ggplot(data = g, aes(y = conc, x = status, fill = as.factor(statCode))) +  
  geom_boxplot() + labs(x = "Disease Status / Classification",  
    y = "Hydroxychloroquine Concentration ( $\mu$  g/ml)",  
    title = "Dose Concentration Plotted Against Disease Status/Classification") +  
  scale_fill_discrete(name = 'Status', labels = c("Asymptomatic",  
    "LRTI",  
    "URTI"))
```